Graph Decomposition*

Xiaofeng Gao

Department of Computer Science and Engineering Shanghai Jiao Tong University, P.R.China

Algorithm Course @ Shanghai Jiao Tong University

Algorithm@SJTU Xiaofeng Gao Graph Decomposition 1/26

Outline

- Depth-First Search in Undirected Graphs
 - Exploring Graphs
 - Connectivity in Undirected Graphs
 - Previsit and Postvisit Orderings
- Depth-First Search in Directed Graphs
 - Types of Edges
 - Directed Acyclic Graphs
 - Strongly Connected Components
- Breadth-First Search
 - Correctness and Efficiency

Exploring Graphs
Connectivity in Undirected Graphs
Previsit and Postvisit Orderings

Exploring Graphs



Exploring Graphs

```
Algorithm 1: EXPLORE(G, v)
  Input: G = (V, E) is a graph; v \in V
  Output: VISITED(u) = true for all nodes u reachable from v
 VISITED(v) = true;
  PREVISIT(v);
 foreach edge(v, u) \in E do
     if not VISITED(u) then
         EXPLORE(G, u);
6 POSTVISIT(v);
```

5

Exploring Graphs

```
Algorithm 1: EXPLORE(G, v)
Input: G = (V, E) is a graph; v \in V
Output: VISITED(u) = true for all nodes u reachable from v

1 VISITED(v) = true;
2 PREVISIT(v);
3 foreach edge(v, u) \in E do
```

6 POSTVISIT(v);

if *not* VISITED(u) then | EXPLORE(G, u);

- PREVISIT, POSTVISIT procedures are optional.
- work on a vertex when first discovered and left for the last time.

Correctness Proof

Theorem: EXPLORE(G, v) is correct (it visits all nodes reachable from v).

4/26

Correctness Proof

Theorem: EXPLORE(G, v) is **correct** (it visits all nodes reachable from v).

Proof: Every node it visits must be reachable from *v*:

EXPLORE moves from node to their neighbors; it can never jump to a region not reachable from v.

Correctness Proof

Theorem: EXPLORE(G, v) is **correct** (it visits all nodes reachable from v).

Proof: Every node it visits must be reachable from *v*:

EXPLORE moves from node to their neighbors; it can never jump to a region not reachable from *v*.

Every node reachable from *v* must be visited:

If $\exists u$ that EXPLORE misses, choose a path from v to u. Let z be the last vertex on that path that EXPLORE visited. Let w be the node immediately after it on this path.

So z was visited but w was not. This is a contradiction: while EXPLORE was at node z, it would have noticed w and moved on to it.

Exploring Graphs
Connectivity in Undirected Graphs
Previsit and Postvisit Orderings

Depth-First Search



Depth-First Search

```
Algorithm 2: DFS(G)
```

```
Input: G = (V, E) is a graph
```

Output: VISITED(v) is set to true for all nodes $v \in V$

- 1 foreach $v \in V$ do
- 2 | VISITED(v) = false;
- 3 foreach $v \in V$ do
- 4 | **if** not VISITED(v) **then**
- 5 | EXPLORE(G, v);

Exploring Graphs Connectivity in Undirected Graph Previsit and Postvisit Orderings

Running Time of DFS



Because of the VISITED array, each vertex is EXPLORE'd just once.

Because of the VISITED array, each vertex is EXPLORE'd just once.

During the exploration of a vertex, there are the following steps:

Because of the VISITED array, each vertex is EXPLORE'd just once.

During the exploration of a vertex, there are the following steps:

Some fixed amount of work − marking the spot as visited, and the PRE/POSTVISIT.

The total work done in this step is then O(|V|).

Because of the VISITED array, each vertex is EXPLORE'd just once.

During the exploration of a vertex, there are the following steps:

- Some fixed amount of work − marking the spot as visited, and the PRE/POSTVISIT.
 - The total work done in this step is then O(|V|).
- A loop in which adjacent edges are scanned, to see if they lead somewhere new.
 - Over the course of the entire DFS, each edge $(x, y) \in E$ is examined exactly *twice*, once during EXPLORE(G, x) and once during EXPLORE(G, y). The overall time is therefore O(|E|).

Because of the VISITED array, each vertex is EXPLORE'd just once.

During the exploration of a vertex, there are the following steps:

- Some fixed amount of work − marking the spot as visited, and the PRE/POSTVISIT.
 - The total work done in this step is then O(|V|).
- ▶ A loop in which adjacent edges are scanned, to see if they lead somewhere new.
 - Over the course of the entire DFS, each edge $(x, y) \in E$ is examined exactly *twice*, once during EXPLORE(G, x) and once during EXPLORE(G, y). The overall time is therefore O(|E|).

Thus the depth-first search has a running time of O(|V| + |E|).

Outline

- Depth-First Search in Undirected Graphs
 - Exploring Graphs
 - Connectivity in Undirected Graphs
 - Previsit and Postvisit Orderings
- Depth-First Search in Directed Graphs
 - Types of Edges
 - Directed Acyclic Graphs
 - Strongly Connected Components
- Breadth-First Search
 - Correctness and Efficiency

When EXPLORE starts at vertex v, it identifies the connected component containing v.



When EXPLORE starts at vertex v, it identifies the connected component containing v.

Each time the DFS outer loop calls EXPLORE, a new connected component is picked out \Rightarrow can check if G is connected.

When EXPLORE starts at vertex v, it identifies the connected component containing v.

Each time the DFS outer loop calls EXPLORE, a new connected component is picked out \Rightarrow can check if G is connected.

More generally, assign each node v an integer CCNUM[v] to identify the connected component to which it belongs.

$$\underline{\mathsf{PREVISIT}}(v)$$

$$CCNUM[v] = cc$$



When EXPLORE starts at vertex v, it identifies the connected component containing v.

Each time the DFS outer loop calls EXPLORE, a new connected component is picked out \Rightarrow can check if G is connected.

More generally, assign each node v an integer CCNUM[v] to identify the connected component to which it belongs.

$$\underline{\text{PREVISIT}}(v)$$

$$CCNUM[v] = cc$$

Initially, cc = 0, will increment each time DFS calls EXPLORE.



Algorithm@SJTU Xiaofeng Gao Graph Decomposition 8/26

Outline

- Depth-First Search in Undirected Graphs
 - Exploring Graphs
 - Connectivity in Undirected Graphs
 - Previsit and Postvisit Orderings
- Depth-First Search in Directed Graphs
 - Types of Edges
 - Directed Acyclic Graphs
 - Strongly Connected Components
- Breadth-First Search
 - Correctness and Efficiency

For each node, we will note down the times of two important events:

- ▶ the moment of first discovery (corresponding to PREVISIT);
- ▶ and the moment of final departure (POSTVISIT).



For each node, we will note down the times of two important events:

- ▶ the moment of first discovery (corresponding to PREVISIT);
- ▶ and the moment of final departure (POSTVISIT).

```
\frac{\text{PREVISIT}(v)}{\text{PRE}[v] = \text{clock}}\frac{\text{clock}}{\text{clock}} = \frac{1}{2}
```

For each node, we will note down the times of two important events:

- ▶ the moment of first discovery (corresponding to PREVISIT);
- ▶ and the moment of final departure (POSTVISIT).

```
\frac{\text{PREVISIT}(v)}{\text{PRE}[v] = \text{clock}} \\
\text{clock} = \text{clock} + 1

\frac{\text{POSTVISIT}(v)}{\text{POST}[v] = \text{clock}} \\
\text{clock} = \text{clock} + 1
```

For each node, we will note down the times of two important events:

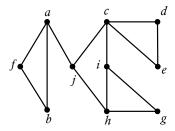
- ▶ the moment of first discovery (corresponding to PREVISIT);
- ▶ and the moment of final departure (POSTVISIT).

```
\frac{\text{PREVISIT}}{\text{PRE}[v]} = \text{clock}
\text{clock} = \text{clock} + 1
\frac{\text{POSTVISIT}}{\text{POST}[v]} = \text{clock}
\text{clock} = \text{clock} + 1
```

Lemma: $\forall u, v \in V$, intervals [PRE(u), POST(u)], [PRE(v), POST(v)] are either *disjoint* or *one is contained within the other*.

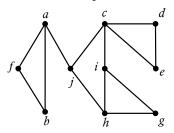
An executing example

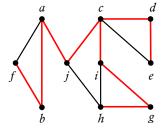
Assume we use alphabetical order to explore *G*:



An executing example

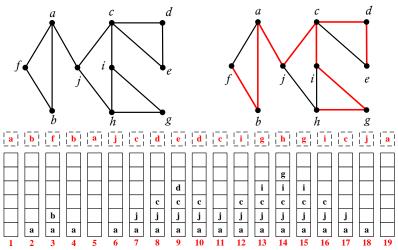
Assume we use alphabetical order to explore *G*:





An executing example

Assume we use alphabetical order to explore *G*:



Outline

- Depth-First Search in Undirected Graphs
 - Exploring Graphs
 - Connectivity in Undirected Graphs
 - Previsit and Postvisit Orderings
- Depth-First Search in Directed Graphs
 - Types of Edges
 - Directed Acyclic Graphs
 - Strongly Connected Components
- Breadth-First Search
 - Correctness and Efficiency

Types of Edges

DFS yields a **search tree/forests**: root; parent and child; descendant and ancestor.



Types of Edges

DFS yields a **search tree/forests**: root; parent and child; descendant and ancestor.

- Tree edges: part of the DFS forest.
- Forward edges: lead from a node to a nonchild descendant in the DFS tree.
- Backedges: lead to an ancestor in the DFS tree.
- **Cross edges**: neither descendant nor ancestor; they lead to a node that has already been explored (that is, already postvisited).

Types of Edges

DFS yields a **search tree/forests**: root; parent and child; descendant and ancestor.

- Tree edges: part of the DFS forest.
- **Forward edges**: lead from a node to a nonchild descendant in the DFS tree.
- Backedges: lead to an ancestor in the DFS tree.
- Cross edges: neither descendant nor ancestor; they lead to a node that has already been explored (that is, already postvisited).

PRE/POST ordering for (u, v)				Edge type
[<i>u</i>	[v]	$]_{v}$	$]_u$	Tree/forward
[v	[u]	$]_u$	$]_v$	Back
[v	$]_{v}$	[u	$]_u$	Cross

Outline

- Depth-First Search in Undirected Graphs
 - Exploring Graphs
 - Connectivity in Undirected Graphs
 - Previsit and Postvisit Orderings
- Depth-First Search in Directed Graphs
 - Types of Edges
 - Directed Acyclic Graphs
 - Strongly Connected Components
- Breadth-First Search
 - Correctness and Efficiency

Types of Edges
Directed Acyclic Graphs
Strongly Connected Components

Directed Acyclic Graphs (DAG)



Algorithm@SJTU Xiaofeng Gao Graph Decomposition 15/26

Directed Acyclic Graphs (DAG)

Lemma: A directed graph has a cycle if and only if its depth-first search reveals a back edge.



Algorithm@SJTU Xiaofeng Gao Graph Decomposition 15/26

Lemma: A directed graph has a cycle if and only if its depth-first search reveals a back edge.

Proof: " \Leftarrow " (easy) If (u, v) is a back edge, then \exists a cycle consisting of this edge together with the path from v to u in the search tree.

Lemma: A directed graph has a cycle if and only if its depth-first search reveals a back edge.

Proof: " \Leftarrow " (easy) If (u, v) is a back edge, then \exists a cycle consisting of this edge together with the path from v to u in the search tree.

" \Rightarrow " Conversely, if the graph has a cycle $v_0 \rightarrow v_2 \rightarrow \cdots v_k \rightarrow v_0$, look at the first node v_i on this cycle to be discovered (the node with the lowest PRE number).

Lemma: A directed graph has a cycle if and only if its depth-first search reveals a back edge.

Proof: " \Leftarrow " (easy) If (u, v) is a back edge, then \exists a cycle consisting of this edge together with the path from v to u in the search tree.

" \Rightarrow " Conversely, if the graph has a cycle $v_0 \rightarrow v_2 \rightarrow \cdots v_k \rightarrow v_0$, look at the first node v_i on this cycle to be discovered (the node with the lowest PRE number).

All the other v_j on the cycle are reachable from it and will therefore be its descendants in the search tree.

Lemma: A directed graph has a cycle if and only if its depth-first search reveals a back edge.

Proof: " \Leftarrow " (easy) If (u, v) is a back edge, then \exists a cycle consisting of this edge together with the path from v to u in the search tree.

" \Rightarrow " Conversely, if the graph has a cycle $v_0 \rightarrow v_2 \rightarrow \cdots v_k \rightarrow v_0$, look at the first node v_i on this cycle to be discovered (the node with the lowest PRE number).

All the other v_j on the cycle are reachable from it and will therefore be its descendants in the search tree.

In particular, the edge $v_{i-1} \rightarrow v_i$ (or $v_k \rightarrow v_0$ if i = 0) is a back edge.



Objective: Order the vertices such that every edge goes from a small vertex to a large one.



Objective: Order the vertices such that every edge goes from a small vertex to a large one.

Lemma: In a dag, every edge leads to a vertex with a lower POST number.

Objective: Order the vertices such that every edge goes from a small vertex to a large one.

Lemma: In a dag, every edge leads to a vertex with a lower POST number.

Hence, a dag can be linearized by decreasing POST numbers, the vertex with the smallest POST number comes last in this linearization, and it must be a **sink** – no outgoing edges. Symmetrically, the one with the highest POST is a **source**, a node with no incoming edges.

Objective: Order the vertices such that every edge goes from a small vertex to a large one.

Lemma: In a dag, every edge leads to a vertex with a lower POST number.

Hence, a dag can be linearized by decreasing POST numbers, the vertex with the smallest POST number comes last in this linearization, and it must be a **sink** – no outgoing edges. Symmetrically, the one with the highest POST is a **source**, a node with no incoming edges.

Lemma: Every dag has at least one source and at least one sink.

Objective: Order the vertices such that every edge goes from a small vertex to a large one.

Lemma: In a dag, every edge leads to a vertex with a lower POST number.

Hence, a dag can be linearized by decreasing POST numbers, the vertex with the smallest POST number comes last in this linearization, and it must be a **sink** – no outgoing edges. Symmetrically, the one with the highest POST is a **source**, a node with no incoming edges.

Lemma: Every dag has at least one source and at least one sink.

The guaranteed existence of a source suggests an alternative approach to linearization:

- ① Find a source, output it, and delete it from the graph.
- 2 Repeat until the graph is empty.



Outline

- Depth-First Search in Undirected Graphs
 - Exploring Graphs
 - Connectivity in Undirected Graphs
 - Previsit and Postvisit Orderings
- Depth-First Search in Directed Graphs
 - Types of Edges
 - Directed Acyclic Graphs
 - Strongly Connected Components
- Breadth-First Search
 - Correctness and Efficiency

Types of Edges Directed Acyclic Graphs Strongly Connected Components

Connectivity for Directed Graphs



Connectivity for Directed Graphs

Definition: Two nodes u and v of a directed graph are **connected** if there is a path from u to v and a path from v to u.



Connectivity for Directed Graphs

Definition: Two nodes u and v of a directed graph are **connected** if there is a path from u to v and a path from v to u.

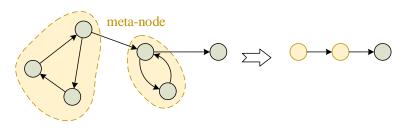
This relation partitions V into disjoint sets that we call **strongly** connected components.

Connectivity for Directed Graphs

Definition: Two nodes u and v of a directed graph are **connected** if there is a path from u to v and a path from v to u.

This relation partitions V into disjoint sets that we call **strongly** connected components.

Lemma: Every directed graph is a dag of its strongly connected components.



Types of Edges
Directed Acyclic Graphs
Strongly Connected Components

Investigation

Lemma: If the EXPLORE subroutine is started at node u, then it will terminate precisely when all nodes reachable from u have been visited.

Lemma: If the EXPLORE subroutine is started at node u, then it will terminate precisely when all nodes reachable from u have been visited.

Therefore, if we call EXPLORE on a node that lies in a sink strongly connected component (a strongly connected component that is a sink in the meta-graph), then we will retrieve exactly that component.

Lemma: If the EXPLORE subroutine is started at node u, then it will terminate precisely when all nodes reachable from u have been visited.

Therefore, if we call EXPLORE on a node that lies in a sink strongly connected component (a strongly connected component that is a sink in the meta-graph), then we will retrieve exactly that component.

Lemma: If C and C' are strongly connected components, and there is an edge from a node in C to a node in C', then the highest POST number in C is bigger than the highest POST number in C'.

Lemma: If the EXPLORE subroutine is started at node u, then it will terminate precisely when all nodes reachable from u have been visited.

Therefore, if we call EXPLORE on a node that lies in a sink strongly connected component (a strongly connected component that is a sink in the meta-graph), then we will retrieve exactly that component.

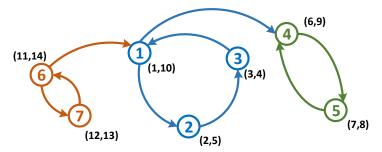
Lemma: If C and C' are strongly connected components, and there is an edge from a node in C to a node in C', then the highest POST number in C is bigger than the highest POST number in C'.

Lemma: The node that receives the highest POST number in a depth-first search must lie in a *source strongly connected component*.

Investigation (Cont')

Note: The smallest POST number in a depth-first search may NOT lie in a *sink strongly connected component!*

An Counter Example: (Node ID denotes the explore order)



The smallest POST number is Node 3, NOT in the sink strongly connected component (green).

Types of Edges Directed Acyclic Graphs Strongly Connected Components

An Efficient Algorithm



To design a linear-time algorithm, we have two problems:

- (A) How do we find a node that we know for sure lies in a sink strongly connected component?
- (B) How do we continue once this first component has been discovered?

To design a linear-time algorithm, we have two problems:

- (A) How do we find a node that we know for sure lies in a sink strongly connected component?
- (B) How do we continue once this first component has been discovered?

Solving Problem A:

Consider the **reverse graph** G^R , the same as G but with all edges reversed (has exactly the same strongly connected components as G).

To design a linear-time algorithm, we have two problems:

- (A) How do we find a node that we know for sure lies in a sink strongly connected component?
- (B) How do we continue once this first component has been discovered?

Solving Problem A:

Consider the **reverse graph** G^R , the same as G but with all edges reversed (has exactly the same strongly connected components as G).

So, if we do a depth-first search of G^R , the node with the highest POST number will come from a source strongly connected component in G^R , which is a sink strongly connected component in G.

Solving Problem B:

Once we have found the first strongly connected component and deleted it from the graph, the node with the highest POST number among those remaining will belong to a sink strongly connected component of whatever remains of *G*.

Solving Problem B:

Once we have found the first strongly connected component and deleted it from the graph, the node with the highest POST number among those remaining will belong to a sink strongly connected component of whatever remains of *G*.

Thus we can keep using the post numbering from our initial depth-first search on G^R to successively output the second strongly connected component, the third strongly connected component, and so on.

Solving Problem B:

Once we have found the first strongly connected component and deleted it from the graph, the node with the highest POST number among those remaining will belong to a sink strongly connected component of whatever remains of *G*.

Thus we can keep using the post numbering from our initial depth-first search on G^R to successively output the second strongly connected component, the third strongly connected component, and so on.

The Linear-Time Algorithm:

Algorithm@SJTU

- ① Run depth-first search on G^R .
- 2 Run depth-first search on G, and process the vertices in decreasing order of their POST numbers from step 1.

Xiaofeng Gao

Graph Decomposition

Outline

- Depth-First Search in Undirected Graphs
 - Exploring Graphs
 - Connectivity in Undirected Graphs
 - Previsit and Postvisit Orderings
- Depth-First Search in Directed Graphs
 - Types of Edges
 - Directed Acyclic Graphs
 - Strongly Connected Components
- Breadth-First Search
 - Correctness and Efficiency

Breadth-First Search



Breadth-First Search

6

7

8

9

10

```
Algorithm 3: BFS(G, s)
  Input: Graph G = (V, E), directed or undirected; vertex s \in V
  Output: DIST(u) is set to the distance from s to all reachable u
1 foreach u \in V do
  DIST(u) = \infty;
3 DIST(s) = 0;
4 Q = [s] (queue containing just s);
5 while Q is not empty do
     u = EJECT(Q);
     foreach edge(u, v) \in E do
         if DIST(v) = \infty then
             INJECT(Q, v);
             DIST(v) = DIST(u) + 1;
```

Correctness and efficiency



25/26

Correctness and efficiency

Lemma: For each d = 0, 1, 2, ..., there is a moment at which

- (1) all nodes at distance $\leq d$ from s have their distances correctly set;
- (2) all other nodes have their distances set to ∞ ; and
- (3) the queue contains exactly the nodes at distance d.



Correctness and efficiency

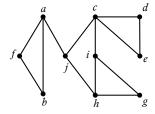
Lemma: For each d = 0, 1, 2, ..., there is a moment at which

- (1) all nodes at distance $\leq d$ from s have their distances correctly set;
- (2) all other nodes have their distances set to ∞ ; and
- (3) the queue contains exactly the nodes at distance d.

Lemma: BFS has a running time of O(|V| + |E|).

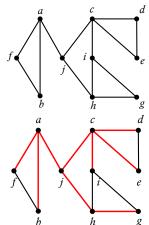
An executing example

Assume we use alphabetical order to explore *G*:



An executing example

Assume we use alphabetical order to explore *G*:



An executing example

Assume we use alphabetical order to explore *G*:

